

▼ Vision Dataset (Pneumonia classification): LTI

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
from google.colab import files
files.upload()
```

```
#create a kaggle folder
! mkdir ~/.kaggle
#copy the Kaggle.json file into Kaggle folder
!cp kaggle.json ~/.kaggle/
```

```
#permission for json to act
! chmod 600 ~/.kaggle/kaggle.json
```

```
! kaggle datasets download -d paultimothymooney/chest-xray-pneumonia
```

Choose files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to e

Saving kaggle.json to kaggle.json

Downloading chest-xray-pneumonia.zip to /content

100% 2.29G/2.29G [00:28<00:00, 40.5MB/s]

100% 2.29G/2.29G [00:28<00:00, 85.1MB/s]

```
! unzip chest-xray-pneumonia.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: chest_xray/train/NORMAL/IM-0435-0001-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0435-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0437-0001-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0437-0001-0002.jpeg
inflating: chest_xray/train/NORMAL/IM-0437-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0438-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0439-0001-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0439-0001-0002.jpeg
inflating: chest_xray/train/NORMAL/IM-0439-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0440-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0441-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0442-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0444-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0445-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0446-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0447-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0448-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0449-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0450-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0451-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0452-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0453-0001-0002.jpeg
inflating: chest_xray/train/NORMAL/IM-0453-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0455-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0456-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0457-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0458-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0459-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0460-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0461-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0463-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0464-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0465-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0466-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0467-0001-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0467-0001-0002.jpeg
inflating: chest_xray/train/NORMAL/IM-0467-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0469-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0471-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0472-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0473-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0474-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0475-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0476-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0477-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0478-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0479-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0480-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0481-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0482-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0483-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0484-0001.jpeg
```

```

inflating: chest_xray/train/NORMAL/IM-0485-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0486-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0487-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0488-0001.jpeg
inflating: chest_xray/train/NORMAL/IM-0489-0001.jpeg

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import tensorflow as tf
from tensorflow.keras import models
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from keras.callbacks import ReduceLROnPlateau
import cv2
import pathlib
# Input data files are available in the "../input/" directory.
import os
for dirname, _, filenames in os.walk('/kaggle/datasets/download -d paultimothymooney/chest-xray-pneumonia'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

#to handle .DS_Store 'hidden files'
train = 'chest_xray/chest_xray/train/NORMAL'
train_P = 'chest_xray/chest_xray/train/PNEUMONIA'
test = 'chest_xray/chest_xray/test/NORMAL'
test_P = 'chest_xray/chest_xray/test/PNEUMONIA'
val = 'chest_xray/chest_xray/val/NORMAL'
val_P = 'chest_xray/chest_xray/val/PNEUMONIA'

train = pathlib.Path(train)
train_P = pathlib.Path(train_P)
test = pathlib.Path(test)
test_P = pathlib.Path(test_P)
val = pathlib.Path(val)
val_P = pathlib.Path(val_P)

# for item in train.glob("*"):
#     print(item.name)

list_ds_N = tf.data.Dataset.list_files(str(train/'*.jpeg'))
list_ds_P = tf.data.Dataset.list_files(str(train_P/'*.jpeg'))

list_ds_test_N = tf.data.Dataset.list_files(str(test/'*.jpeg'))
list_ds_test_P = tf.data.Dataset.list_files(str(test_P/'*.jpeg'))

list_ds_val_N = tf.data.Dataset.list_files(str(val/'*.jpeg'))
list_ds_val_P = tf.data.Dataset.list_files(str(val_P/'*.jpeg'))

for f in list_ds_P.take(5):
    print(f.numpy())

    b'chest_xray/chest_xray/train/PNEUMONIA/person758_bacteria_2662.jpeg'
    b'chest_xray/chest_xray/train/PNEUMONIA/person1644_bacteria_4357.jpeg'
    b'chest_xray/chest_xray/train/PNEUMONIA/person803_bacteria_2710.jpeg'
    b'chest_xray/chest_xray/train/PNEUMONIA/person700_bacteria_2599.jpeg'
    b'chest_xray/chest_xray/train/PNEUMONIA/person292_virus_599.jpeg'

list_ds = list_ds_N.concatenate(list_ds_P)
list_ds_test = list_ds_test_N.concatenate(list_ds_test_P)
list_ds_val = list_ds_val_N.concatenate(list_ds_val_P)

# Reads an image from a file, decodes it into a dense tensor, and resizes it
# to a fixed shape.
def parse_image(filename):
    parts = tf.strings.split(file_path, '/')
    label = parts[-2]

    image = tf.io.read_file(filename)
    image = tf.io.decode_image(image, expand_animations = False)
    image = tf.image.convert_image_dtype(image, tf.float32)
    image = tf.image.resize(image, [150, 150])
    return image, label

def show(image, label):
    plt.figure()

```

```
plt.imshow( tf.squeeze(image))
plt.title(label.numpy().decode('utf-8'))
plt.axis('off')
```

▼ Data Visualization and augmentation



Plot at least two samples from each class of the dataset

▼ Plotting sample data from training data set

```
print('Plotting sample dataset from Training data')
```

```
file_path = next(iter(list_ds))
image, label = parse_image(file_path)
show(image, label)
print("-----")
file_path = next(iter(list_ds))
image, label = parse_image(file_path)
show(image, label)
```

```
Plotting sample dataset from Training data
-----
```




▼ Plotting sample data from testing data set

```
print('Plotting sample dataset from Testing data')
```

```
file_path = next(iter(list_ds_test))
image, label = parse_image(file_path)
show(image, label)
print("-----")
file_path = next(iter(list_ds_test))
image, label = parse_image(file_path)
show(image, label)
```

Ploting sample dataset from Testing data

NORMAL



▼ Ploting sample data from validation data set




```
print('Ploting sample dataset from validation data')
```


```
file_path = next(iter(list_ds_val))
image, label = parse_image(file_path)
show(image, label)
print("-----")
file_path = next(iter(list_ds_val))
image, label = parse_image(file_path)
show(image, label)
```

Ploting sample dataset from validation data

NORMAL



NORMAL



Bringing the train and test data in the required format.

```
# Reading and Parsing the dataset
images_ds = list_ds.map(parse_image)
images_ds_test = list_ds_test.map(parse_image)
images_ds_val = list_ds_val.map(parse_image)
```

▼ Apply rotation and height shift augmentatio

```
# defining rotate and shift function

import scipy.ndimage as ndimage
def random_rotate_image(image):
    image = ndimage.rotate(image, np.random.uniform(-30,30), reshape=False)
    return image

def random_shift_image(image):
    image = ndimage.shift(image, shift = 0.080)
    return image
```

Ploting sample data post rotation and height shift on training data

```

print('Ploting sample dataset from Training data post rotation')

image, label = next(iter(images_ds))
image = random_rotate_image(image)
show(image, label)
print('Ploting sample dataset from Training data')
show(image, label)

print("-----")
print('Ploting sample dataset from Training data post rotation')

image, label = next(iter(images_ds))
image = random_rotate_image(image)
show(image, label)

print('Ploting sample dataset from Training data')
show(image, label)


```

```


Ploting sample dataset from Training data post rotation
Ploting sample dataset from Training data
-----
Ploting sample dataset from Training data post rotation
Ploting sample dataset from Training data

```


NORMAL




NORMAL



NORMAL



NORMAL



```

print('Ploting sample dataset from Testing data post rotation')

```

```

image, label = next(iter(images_ds_test))
image = random_rotate_image(image)
show(image, label)
print('Ploting sample dataset from test data')
show(image, label)

print("-----")
print('Ploting sample dataset from Testing data post rotation')

image, label = next(iter(images_ds_test))
image = random_rotate_image(image)
show(image, label)


print('Ploting sample dataset from test data')
show(image, label)

```


Ploting sample dataset from Testing data post rotation
 Ploting sample dataset from test data

 Ploting sample dataset from Testing data post rotation
 Ploting sample dataset from test data


NORMAL




NORMAL



NORMAL



NORMAL



```

print('Ploting sample dataset from validation data post rotation')

image, label = next(iter(images_ds_val))
image = random_rotate_image(image)
show(image, label)
print('Ploting sample dataset from validation data')
show(image, label)

```

```

print("-----")
print('Ploting sample dataset from validation data post rotation')

image, label = next(iter(images_ds_val))
image = random_rotate_image(image)
show(image, label)

print('Ploting sample dataset from validation data')
show(image, label)


```

```


Ploting sample dataset from validation data post rotation
Ploting sample dataset from validation data
-----
Ploting sample dataset from validation data post rotation
Ploting sample dataset from validation data

```


NORMAL




NORMAL



NORMAL



NORMAL



```

# def tf_random_rotate_image(image, label):
#     im_shape = image.shape
#     [image,] = tf.py_function(random_rotate_image, [image], [tf.float32])
#     image.set_shape(im_shape)
#     return image, label

# rot_ds = images_ds.map(tf_random_rotate_image)
# rot_ds_test = images_ds_test.map(tf_random_rotate_image)
# rot_ds_val = images_ds_val.map(tf_random_rotate_image)

```

```
# image, label = next(iter(rot_ds_val))
# plt.figure(figsize = (5,5))
# plt.imshow(tf.squeeze(image), cmap='gray')
# plt.title(label)
```

- ▼ **Print the shapes of train and test data.**

```
image, label = next(iter(images_ds))
print('shape of the training data', len(list_ds), image.shape[0])
print("-----")
image, label = next(iter(images_ds_test))
print('shape of the testing data', len(list_ds_test), image.shape[0])
print("-----")
image, label = next(iter(images_ds_val))
print('shape of the validation data', len(list_ds_val), image.shape[0])
```

```
shape of the training data 5216 150
-----
shape of the testing data 624 150
-----
shape of the validation data 16 150
```

```
labels = ['PNEUMONIA', 'NORMAL']
img_size = 150
def get_training_data(data_dir):
    data = []
    for label in labels:
        path = os.path.join(data_dir, label)
        class_num = labels.index(label)
        for img in os.listdir(path):
            try:
                img_arr = cv2.imread(os.path.join(path, img), cv2.IMREAD_GRAYSCALE)
                resized_arr = cv2.resize(img_arr, (img_size, img_size)) # Reshaping images to preferred size
                data.append([resized_arr, class_num])
            except Exception as e:
                print(e)
    return np.array(data)
```

```
train = get_training_data('chest_xray/chest_xray/train/')
test = get_training_data('chest_xray/chest_xray/test/')
val = get_training_data('chest_xray/chest_xray/val/')

```

```
OpenCV(4.1.2) /io/opencv/modules/imgproc/src/resize.cpp:3720: error: (-215:Assertion failed) !ssize.empty() in function '
OpenCV(4.1.2) /io/opencv/modules/imgproc/src/resize.cpp:3720: error: (-215:Assertion failed) !ssize.empty() in function '
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:15: VisibleDeprecationWarning: Creating an ndarray from ragged
  from ipykernel import kernelapp as app
OpenCV(4.1.2) /io/opencv/modules/imgproc/src/resize.cpp:3720: error: (-215:Assertion failed) !ssize.empty() in function '
OpenCV(4.1.2) /io/opencv/modules/imgproc/src/resize.cpp:3720: error: (-215:Assertion failed) !ssize.empty() in function '
```

```
x_train = []
y_train = []
```

```
x_val = []
y_val = []
```

```
x_test = []
y_test = []
```

```
for feature, label in train:
    x_train.append(feature)
    y_train.append(label)
```

```
for feature, label in test:
    x_test.append(feature)
    y_test.append(label)
```



```

for feature, label in val:
    x_val.append(feature)
    y_val.append(label)

# Normalize the data
x_train = np.array(x_train) / 255
x_val = np.array(x_val) / 255
x_test = np.array(x_test) / 255

# y_train = np.array(y_train)
# y_val = np.array(y_val)
# y_test = np.array(y_test)
y_train = tf.keras.utils.to_categorical(y_train)
y_test = tf.keras.utils.to_categorical(y_test)
y_val = tf.keras.utils.to_categorical(y_val)

Xtrain = x_train.reshape((5216, 150*150))
Xtest = x_test.reshape((624, 150*150))
Xval = x_val.reshape((16, 150*150))

```

▼ Model building

```

# Create a model object
from keras.layers import Dropout

tf.keras.backend.clear_session() #clear old tf session and models
dnnModel = models.Sequential()
# Layer 1 = input layer
# specify the input size in the first layer.

dnnModel.add(layers.Dense(50, activation='relu', input_shape= (150*150,), kernel_regularizer='l2'))

# Layer 2 = hidden layer
dnnModel.add(layers.Dense(60, activation='relu', kernel_regularizer='l2'))
# dropout in hidden layers with weight constraint
dnnModel.add(Dropout(0.2))
# Layer 3 = hidden layer
dnnModel.add(layers.Dense(30, activation='relu', kernel_regularizer='l2'))

# Layer 4 = output layer
dnnModel.add(layers.Dense(2, activation='softmax', kernel_regularizer='l2'))

dnnModel.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 50)	1125050
dense_1 (Dense)	(None, 60)	3060
dropout (Dropout)	(None, 60)	0
dense_2 (Dense)	(None, 30)	1830
dense_3 (Dense)	(None, 2)	62

```

Total params: 1,130,002
Trainable params: 1,130,002
Non-trainable params: 0

```

▼ Model Compilation

```
dnnModel.compile( optimizer = 'adam', loss = 'binary_crossentropy', metrics=['accuracy', 'mse'] )
```

▼ Model Training

```
import datetime
t1 = datetime.datetime.now()
h = dnnModel.fit( Xtrain, y_train, batch_size=64, validation_data = (Xtest, y_test), epochs=200)
t2 = datetime.datetime.now()
```

Epoch 1/200
82/82 [=====] - 3s 29ms/step - loss: 1.8098 - accuracy: 0.8127 - mse: 0.1363 - val_loss: 1.5892
Epoch 2/200
82/82 [=====] - 2s 26ms/step - loss: 1.0375 - accuracy: 0.8854 - mse: 0.0845 - val_loss: 1.0381
Epoch 3/200
82/82 [=====] - 2s 26ms/step - loss: 0.7184 - accuracy: 0.9218 - mse: 0.0572 - val_loss: 1.4043
Epoch 4/200
82/82 [=====] - 2s 27ms/step - loss: 0.6008 - accuracy: 0.9260 - mse: 0.0547 - val_loss: 1.1106
Epoch 5/200
82/82 [=====] - 4s 43ms/step - loss: 0.5104 - accuracy: 0.9363 - mse: 0.0466 - val_loss: 0.7531
Epoch 6/200
82/82 [=====] - 3s 38ms/step - loss: 0.4986 - accuracy: 0.9277 - mse: 0.0543 - val_loss: 1.2790
Epoch 7/200
82/82 [=====] - 2s 26ms/step - loss: 0.4843 - accuracy: 0.9225 - mse: 0.0571 - val_loss: 1.1906
Epoch 8/200
82/82 [=====] - 2s 27ms/step - loss: 0.4118 - accuracy: 0.9459 - mse: 0.0422 - val_loss: 0.8347
Epoch 9/200
82/82 [=====] - 2s 27ms/step - loss: 0.3971 - accuracy: 0.9421 - mse: 0.0447 - val_loss: 0.7081
Epoch 10/200
82/82 [=====] - 2s 26ms/step - loss: 0.4300 - accuracy: 0.9172 - mse: 0.0609 - val_loss: 0.6804
Epoch 11/200
82/82 [=====] - 2s 28ms/step - loss: 0.3918 - accuracy: 0.9327 - mse: 0.0503 - val_loss: 0.5622
Epoch 12/200
82/82 [=====] - 2s 25ms/step - loss: 0.3839 - accuracy: 0.9317 - mse: 0.0514 - val_loss: 0.5947
Epoch 13/200
82/82 [=====] - 2s 26ms/step - loss: 0.3582 - accuracy: 0.9383 - mse: 0.0479 - val_loss: 1.0117
Epoch 14/200
82/82 [=====] - 2s 27ms/step - loss: 0.3313 - accuracy: 0.9434 - mse: 0.0431 - val_loss: 1.2462
Epoch 15/200
82/82 [=====] - 2s 24ms/step - loss: 0.3463 - accuracy: 0.9344 - mse: 0.0500 - val_loss: 0.7679
Epoch 16/200
82/82 [=====] - 2s 27ms/step - loss: 0.3237 - accuracy: 0.9369 - mse: 0.0472 - val_loss: 0.8396
Epoch 17/200
82/82 [=====] - 2s 25ms/step - loss: 0.3030 - accuracy: 0.9465 - mse: 0.0417 - val_loss: 0.9019
Epoch 18/200
82/82 [=====] - 2s 26ms/step - loss: 0.2930 - accuracy: 0.9408 - mse: 0.0436 - val_loss: 0.8371
Epoch 19/200
82/82 [=====] - 2s 26ms/step - loss: 0.3135 - accuracy: 0.9293 - mse: 0.0505 - val_loss: 0.7632
Epoch 20/200
82/82 [=====] - 2s 26ms/step - loss: 0.2828 - accuracy: 0.9438 - mse: 0.0426 - val_loss: 0.6930
Epoch 21/200
82/82 [=====] - 2s 26ms/step - loss: 0.2593 - accuracy: 0.9492 - mse: 0.0381 - val_loss: 1.3076
Epoch 22/200
82/82 [=====] - 2s 25ms/step - loss: 0.2734 - accuracy: 0.9402 - mse: 0.0440 - val_loss: 1.1349
Epoch 23/200
82/82 [=====] - 2s 26ms/step - loss: 0.2576 - accuracy: 0.9461 - mse: 0.0404 - val_loss: 0.8486
Epoch 24/200
82/82 [=====] - 2s 25ms/step - loss: 0.2480 - accuracy: 0.9503 - mse: 0.0382 - val_loss: 0.6571
Epoch 25/200
82/82 [=====] - 2s 26ms/step - loss: 0.2430 - accuracy: 0.9500 - mse: 0.0390 - val_loss: 1.7212
Epoch 26/200
82/82 [=====] - 2s 25ms/step - loss: 0.2386 - accuracy: 0.9498 - mse: 0.0390 - val_loss: 1.2329
Epoch 27/200
82/82 [=====] - 2s 24ms/step - loss: 0.2618 - accuracy: 0.9398 - mse: 0.0462 - val_loss: 0.7149
Epoch 28/200
82/82 [=====] - 1s 18ms/step - loss: 0.2331 - accuracy: 0.9486 - mse: 0.0388 - val_loss: 0.6511
Epoch 29/200
82/82 [=====] - 1s 18ms/step - loss: 0.2281 - accuracy: 0.9532 - mse: 0.0376 - val_loss: 0.8447

▼ summarize history for accuracy and loss


```
print('time taken by the model', t2-t1)

time taken by the model 0:06:26.746463

# summarize history for accuracy
def plot_acc(history):
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()
# summarize history for loss
def plot_loss(history):
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
```

```
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
plot_acc(h)
plot_loss(h)
```



Model Evaluation

```
#prediction
result = dnnModel.predict(Xval)
print(result.shape)

result = pd.DataFrame(result)
result.columns = ['probal', 'proba2']
result['output'] = result['probal'].apply(lambda x : 1 if x >= 0.50 else 0)
result.head()
```

WARNING:tensorflow:5 out of the last 13 calls to <function Model.make_predict_ (16, 2)

	probal	proba2	output
0	0.998612	0.001388	1
1	0.996382	0.003618	1
2	0.997664	0.002336	1
3	0.995824	0.004176	1
4	0.998222	0.001778	1

```
yval = pd.DataFrame(y_val)
yval.columns = ['trueval1', 'trueval2']
yval = yval[['trueval1']]
```

```
yval.head()
```

truevall

Model confusion matrix report


```
import seaborn as sns
from sklearn.metrics import confusion_matrix
cf_matrix = confusion_matrix(yval['truevall'], result['output'])

ax = sns.heatmap(cf_matrix, annot=True, cmap='Blues')

ax.set_title('Seaborn Confusion Matrix with labels\n\n');
ax.set_xlabel('\nPredicted Values');
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



Model classification report

```
from sklearn.metrics import classification_report

target_names = ['class pneumonia', 'class normal']
print(classification_report(yval['truevall'], result['output'], target_names=target_names))
```

	precision	recall	f1-score	support
class pneumonia	1.00	0.12	0.22	8
class normal	0.53	1.00	0.70	8
accuracy			0.56	16
macro avg	0.77	0.56	0.46	16
weighted avg	0.77	0.56	0.46	16

Hyperparameter Tuning

1. Optimiser: Use a different optimizer with the appropriate LR value.

```
# learning rate schedule
def step_decay(epoch):
    initial_lrate = 0.1
    drop = 0.5
    epochs_drop = 10.0
    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
    return lrate
```

```
# Create a model object
from keras.layers import Dropout
from keras.callbacks import LearningRateScheduler
import math
import random
import numpy as np

tf.keras.backend.clear_session() #clear old tf seesion and models
dnnModel = models.Sequential()
# Layer 1 = input layer
# specify the input size in the first layer.

dnnModel.add(layers.Dense(50, activation='relu', input_shape= (150*150,), kernel_regularizer='l2'))

# Layer 2 = hidden layer
dnnModel.add(layers.Dense(60, activation='relu', kernel_regularizer='l2'))
# dropout in hidden layers with weight constraint
dnnModel.add(Dropout(0.2))
# Layer 3 = hidden layer
dnnModel.add(layers.Dense(30, activation='relu', kernel_regularizer='l2'))

# Layer 4 = output layer
dnnModel.add(layers.Dense(2, activation='softmax', kernel_regularizer='l2'))

dnnModel.summary()
# learning schedule callback
lrate = LearningRateScheduler(step_decay)
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=50)

# Configure the model for training, by using appropriate optimizers and regularizations
# Available optimizer: adam, rmsprop, adagrad, sgd
# loss: objective that the model will try to minimize.
# Available loss: categorical_crossentropy, binary_crossentropy, mean_squared_error
# metrics: List of metrics to be evaluated by the model during training and testing.
dnnModel.compile(optimizer = tf.keras.optimizers.SGD(learning_rate=0.0, momentum=0.9), loss = 'binary_crossentropy', metrics=[
h = dnnModel.fit( Xtrain, y_train, batch_size=64, validation_data = (Xtest, y_test), epochs=50, callbacks=[callback, lrate])
```

```

Epoch 41/50
82/82 [=====] - 1s 17ms/step - loss: 0.5701 - accuracy: 0.7429 - mse: 0.2132 - val_loss: 0.6951
Epoch 48/50
82/82 [=====] - 1s 17ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2136 - val_loss: 0.6954
Epoch 49/50
82/82 [=====] - 1s 17ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2137 - val_loss: 0.6952
Epoch 50/50
82/82 [=====] - 1s 17ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2134 - val_loss: 0.6951

```

```

#prediction
result = dnnModel.predict(Xval)
print(result.shape)

result = pd.DataFrame(result)
result.columns = ['probal', 'proba2']
result['output'] = result['probal'].apply(lambda x : 1 if x >= 0.50 else 0)
result.head()

```

```

(16, 2)

```

	probal	proba2	output
0	0.892996	0.107004	1
1	0.892996	0.107004	1
2	0.892996	0.107004	1
3	0.892996	0.107004	1
4	0.892996	0.107004	1

```

import seaborn as sns
from sklearn.metrics import confusion_matrix
cf_matrix = confusion_matrix(yval['trueval1'], result['output'])

```

```
ax = sns.heatmap(cf_matrix, annot=True, cmap='Blues')
```

```

ax.set_title('Seaborn Confusion Matrix with labels\n\n');
ax.set_xlabel('\nPredicted Values');
ax.set_ylabel('Actual Values ');

```

```

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])


```

```

## Display the visualization of the Confusion Matrix.
plt.show()

```

Seaborn Confusion Matrix with labels



```

from sklearn.metrics import classification_report

target_names = ['class pneumonia', 'class normal']
print(classification_report(yval['trueval1'], result['output'], target_names=target_names))

```

	precision	recall	f1-score	support
class pneumonia	0.00	0.00	0.00	8
class normal	0.50	1.00	0.67	8
accuracy			0.50	16
macro avg	0.25	0.50	0.33	16
weighted avg	0.25	0.50	0.33	16

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to NaN due to the lack of true positives:
warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to NaN due to the lack of true positives:
warn_prf(average, modifier, msg_start, len(result))

```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score  
_warn_prf(average, modifier, msg_start, len(result))
```

```
# Create a model object
from keras.layers import Dropout
from keras.callbacks import LearningRateScheduler
import math
import random
import numpy as np

tf.keras.backend.clear_session() #clear old tf session and models
dnnModel = models.Sequential()
# Layer 1 = input layer
# specify the input size in the first layer.

dnnModel.add(layers.Dense(50, activation='relu', input_shape= (150*150,), kernel_regularizer='l2'))
# Layer 2 = hidden layer
dnnModel.add(layers.Dense(80, activation='relu', kernel_regularizer='l2'))
# dropout in hidden layers with weight constraint
dnnModel.add(Dropout(0.2))
# Layer 3 = hidden layer
dnnModel.add(layers.Dense(60, activation='relu', kernel_regularizer='l2'))
# dropout in hidden layers with weight constraint
dnnModel.add(Dropout(0.2))
# Layer 4 = hidden layer
dnnModel.add(layers.Dense(30, activation='relu', kernel_regularizer='l2'))
# Layer 5 = hidden layer
dnnModel.add(layers.Dense(10, activation='relu', kernel_regularizer='l2'))
# Layer 6 = output layer
dnnModel.add(layers.Dense(2, activation='softmax', kernel_regularizer='l2'))

dnnModel.summary()
# learning schedule callback
lr = LearningRateScheduler(step_decay)
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=50)

# Configure the model for training, by using appropriate optimizers and regularizations
# Available optimizer: adam, rmsprop, adagrad, SGD
# loss: objective that the model will try to minimize.
# Available loss: categorical_crossentropy, binary_crossentropy, mean_squared_error
# metrics: List of metrics to be evaluated by the model during training and testing.
dnnModel.compile(optimizer = tf.keras.optimizers.SGD(learning_rate=0.0, momentum=0.9), loss = 'binary_crossentropy', metrics=[
h = dnnModel.fit( Xtrain, y_train, batch_size=64, validation_data = (Xtest, y_test), epochs=50, callbacks=[callback, lr])
```

```

82/82 [=====] - 3s 39ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2134 - val_loss: 0.6947
Epoch 43/50
82/82 [=====] - 2s 26ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2134 - val_loss: 0.6948
Epoch 44/50
82/82 [=====] - 2s 24ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2132 - val_loss: 0.6948
Epoch 45/50
82/82 [=====] - 2s 26ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2137 - val_loss: 0.6950
Epoch 46/50
82/82 [=====] - 2s 27ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2134 - val_loss: 0.6948
Epoch 47/50
82/82 [=====] - 2s 24ms/step - loss: 0.5701 - accuracy: 0.7429 - mse: 0.2142 - val_loss: 0.6954
Epoch 48/50
82/82 [=====] - 2s 25ms/step - loss: 0.5701 - accuracy: 0.7429 - mse: 0.2134 - val_loss: 0.6953
Epoch 49/50
82/82 [=====] - 2s 23ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2136 - val_loss: 0.6949
Epoch 50/50
82/82 [=====] - 2s 24ms/step - loss: 0.5700 - accuracy: 0.7429 - mse: 0.2132 - val_loss: 0.6947

```

```

#prediction
result = dnnModel.predict(Xval)
print(result.shape)

result = pd.DataFrame(result)
result.columns = ['probal', 'proba2']
result['output'] = result['probal'].apply(lambda x : 1 if x >= 0.50 else 0)
result.head()

```

```

(16, 2)

```

	probal	proba2	output
0	0.892324	0.107676	1
1	0.892324	0.107676	1
2	0.892324	0.107676	1
3	0.892324	0.107676	1
4	0.892324	0.107676	1

```

import seaborn as sns
from sklearn.metrics import confusion_matrix
cf_matrix = confusion_matrix(yval['trueval'], result['output'])

ax = sns.heatmap(cf_matrix, annot=True, cmap='Blues')


ax.set_title('Seaborn Confusion Matrix with labels\n\n');
ax.set_xlabel('\nPredicted Values');
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()

```

Seaborn Confusion Matrix with labels



```

from sklearn.metrics import classification_report

target_names = ['class pneumonia', 'class normal']
print(classification_report(yval['trueval'], result['output'], target_names=target_names))

```

	precision	recall	f1-score	support
class pneumonia	0.00	0.00	0.00	8
class normal	0.50	1.00	0.67	8

accuracy			0.50	16
macro avg	0.25	0.50	0.33	16
weighted avg	0.25	0.50	0.33	16

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score
_warn_prf(average, modifier, msg_start, len(result))

```

▼ Description on function used in model building

Problem Statement: Classification

Output: Binary output

Final activation function: Sigmoid

Loss Function: Binary cross entropy/ categorical_crossentropy

ReLU — This results in a numerical value greater than 0

Sigmoid — This results in a value between 0 and 1 which we can infer to be how confident it is of it being in the class

Binary Cross Entropy — Cross entropy quantifies the difference between two probability distribution. Our model predicts a model distribution of $\{p, 1-p\}$ (binary distribution) for each of the classes. We use binary cross-entropy to compare these with the true distributions $\{y, 1-y\}$ for each class and sum up their results

Conclusion:

Adam optimization has performed better with training accuracy of 95.63% than Stochastic Gradient Descent (SGD) model with accuracy of 74.0%